

1. Counting missing values

Sports clothing and athleisure attire is a huge industry, worth approximately [\\$193 billion in 2021](#) with a strong growth forecast over the next decade!

In this notebook, we play the role of a product analyst for an online sports clothing company. The company is specifically interested in how it can improve revenue. We will dive into product data such as pricing, reviews, descriptions, and ratings, as well as revenue and website traffic, to produce recommendations for its marketing and sales teams.

The database provided to us, `sports`, contains five tables, with `product_id` being the primary key for all of them:

info

column	data type	description
<code>product_name</code>	<code>varchar</code>	Name of the product
<code>product_id</code>	<code>varchar</code>	Unique ID for product
<code>description</code>	<code>varchar</code>	Description of the product

finance

column	data type	description
<code>product_id</code>	<code>varchar</code>	Unique ID for product
<code>listing_price</code>	<code>float</code>	Listing price for product
<code>sale_price</code>	<code>float</code>	Price of the product when on sale
<code>discount</code>	<code>float</code>	Discount, as a decimal, applied to the sale price
<code>revenue</code>	<code>float</code>	Amount of revenue generated by each product, in US dollars

reviews

column	data type	description
<code>product_name</code>	<code>varchar</code>	Name of the product
<code>product_id</code>	<code>varchar</code>	Unique ID for product
<code>rating</code>	<code>float</code>	Product rating, scored from <code>1.0</code> to <code>5.0</code>
<code>reviews</code>	<code>float</code>	Number of reviews for the product

traffic

column	data type	description
product_id	varchar	Unique ID for product
last_visited	timestamp	Date and time the product was last viewed on the website

brands

column	data type	description
product_id	varchar	Unique ID for product
brand	varchar	Brand of the product

We will be dealing with missing data as well as numeric, string, and timestamp data types to draw insights about the products in the online store. Let's start by finding out how complete the data is.

```
In [2]: %%sql
postgres://sports

-- Count all columns as total_rows
-- Count the number of non-missing entries for description, listing_price, and last_v
-- Join info, finance, and traffic

SELECT
    COUNT(i.*) AS total_rows,
    SUM(CASE WHEN i.description IS NULL THEN 0 ELSE 1 END) AS count_description,
    SUM(CASE WHEN f.listing_price IS NULL THEN 0 ELSE 1 END) AS count_listing_price,
    SUM(CASE WHEN t.last_visited IS NULL THEN 0 ELSE 1 END) AS count_last_visited
FROM info AS i
JOIN finance AS f
    ON i.product_id = f.product_id
JOIN traffic AS t
    ON i.product_id = t.product_id
```

1 rows affected.

```
Out[2]: total_rows  count_description  count_listing_price  count_last_visited
        3179             3117             3120             2928
```

```
In [3]: %%nose

last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert "total_rows" in set(last_output_df.columns), \
        ""Did you alias the count of all products as "total_rows"?""
    assert set(last_output_df.columns) == set(['total_rows', 'count_description', 'cou
        'count_last_visited']), \
        ""Did you select four columns and use the aliases in the instructions?""
```

```
def test_shape():
    assert last_output_df.shape[0] == 1, \
        """Did you return a single row containing the count of values for each column?"""
    assert last_output_df.shape[1] == 4, \
        """Did you select four columns?"""

def test_values():
    assert last_output_df.values.tolist() == [[3179, 3117, 3120, 2928]], \
        """Did you correctly calculate the values for each column? Expected different results"""
```

Out[3]: 3/3 tests passed

2. Nike vs Adidas pricing

We can see the database contains 3,179 products in total. Of the columns we previewed, only one — `last_visited` — is missing more than five percent of its values. Now let's turn our attention to pricing.

How do the price points of Nike and Adidas products differ? Answering this question can help us build a picture of the company's stock range and customer market. We will run a query to produce a distribution of the `listing_price` and the count for each price, grouped by `brand`.

```
In [4]: %%sql

-- Select the brand, listing_price as an integer, and a count of all products in finance
-- Join brands to finance on product_id
-- Aggregate results by brand and listing_price, and sort the results by listing_price
-- Filter for products with a listing_price more than zero

SELECT
    b.brand,
    CAST(f.listing_price AS INT),
    COUNT(f.*)
FROM brands AS b
JOIN finance AS f
    ON b.product_id = f.product_id
WHERE CAST(f.listing_price AS INT) > 0
GROUP BY
    b.brand,
    CAST(f.listing_price AS INT)
ORDER BY CAST(f.listing_price AS INT) DESC
;

* postgresql:///sports
77 rows affected.
```

Out[4]:

brand	listing_price	count
-------	---------------	-------

Adidas	300	2
Adidas	280	4
Adidas	240	5
Adidas	230	8
Adidas	220	11
Adidas	200	8
Nike	200	1
Nike	190	2
Adidas	190	7
Adidas	180	34
Nike	180	4
Nike	170	14
Adidas	170	27
Nike	160	31
Adidas	160	28
Adidas	150	41
Nike	150	6
Nike	140	12
Adidas	140	36
Adidas	130	96
Nike	130	12
Nike	120	16
Adidas	120	115
Nike	110	17
Adidas	110	91
Nike	100	14
Adidas	100	72
Adidas	96	2
Nike	95	1
Adidas	90	89
Nike	90	13
Adidas	86	7
Adidas	85	1
Nike	85	5

Nike	80	16
Adidas	80	322
Nike	79	1
Adidas	76	149
Nike	75	7
Adidas	75	1
Nike	70	4
Adidas	70	87
Adidas	66	102
Nike	65	1
Adidas	63	1
Nike	60	2
Adidas	60	211
Adidas	56	174
Adidas	55	2
Adidas	53	43
Adidas	50	183
Nike	50	5
Adidas	48	42
Nike	48	1
Adidas	46	163
Nike	45	3
Adidas	45	1
Adidas	43	51
Adidas	40	81
Nike	40	1
Adidas	38	24
Adidas	36	25
Adidas	33	24
Adidas	30	37
Nike	30	2
Adidas	28	38
Adidas	27	18
Adidas	25	28
Adidas	23	1

Adidas	20	8
Adidas	18	4
Adidas	16	4
Adidas	15	27
Adidas	13	27
Adidas	12	1
Adidas	10	11
Adidas	9	1

```
In [5]: %%nose

last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert set(last_output_df.columns) == set(['brand', 'count', 'listing_price']), \
        ""Did you select the correct columns and alias the first as "total_rows"?""

def test_shape():
    assert last_output_df.shape[0] == 77, \
        ""Did you correctly aggregate by brand? Expected the output to contain 77 product
    assert last_output_df.shape[1] == 3, \
        ""The output should contain three columns: "brand", "listing_price", and "count"?

def test_values():
    assert last_output_df.iloc[0].values.tolist() == ['Adidas', 300, 2], \
        ""Did you sort the results by "listing_price" in descending order?""
```

Out[5]: 3/3 tests passed

3. Labeling price ranges

It turns out there are 77 unique prices for the products in our database, which makes the output of our last query quite difficult to analyze.

Let's build on our previous query by assigning labels to different price ranges, grouping by `brand` and `label`. We will also include the total `revenue` for each price range and `brand`.

```
In [6]: %%sql

-- Select the brand, a count of all products in the finance table, and total revenue
-- Create four labels for products based on their price range, aliasing as price_category
-- Join brands to finance on product_id
-- Group results by brand and price_category, sort by total_revenue and filter out products with zero revenue

SELECT
    b.brand,
    COUNT(f.*),
```

```

SUM(f.revenue) AS total_revenue,

CASE WHEN f.listing_price < 42 THEN 'Budget'
WHEN f.listing_price >= 42 AND f.listing_price < 74 THEN 'Average'
WHEN f.listing_price >= 74 AND f.listing_price < 129 THEN 'Expensive'
ELSE 'Elite' END AS price_category

FROM brands AS b
INNER JOIN finance AS f
    ON b.product_id = f.product_id
GROUP BY
    b.brand,
    price_category
HAVING b.brand IS NOT NULL
ORDER BY total_revenue DESC
;

```

* postgresql:///sports
8 rows affected.

Out[6]: **brand** **count** **total_revenue** **price_category**

Adidas	849	4626980.0699999999	Expensive
Adidas	1060	3233661.0600000001	Average
Adidas	307	3014316.8299999987	Elite
Adidas	359	651661.1200000002	Budget
Nike	357	595341.0199999992	Budget
Nike	82	128475.59000000003	Elite
Nike	90	71843.15000000004	Expensive
Nike	16	6623.5	Average

```

In [7]: %%nose

last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert set(last_output_df.columns) == set(['brand', 'price_category', 'count', 'total_revenue'])
    """Did you select the correct columns? Expected "brand", "listing_price", "count", "total_revenue" """

def test_shape():
    assert last_output_df.shape[0] == 8, \
        "Did you group by brand and labels? Expected there to be eight rows."
    assert last_output_df.shape[1] == 4, \
        "Did you select four columns?"

def test_values():
    assert last_output_df[:4].values.tolist() == [['Adidas', 849, 4626980.0699999999, 'Expensive'],
        ['Adidas', 1060, 3233661.0600000001, 'Average'],
        ['Adidas', 307, 3014316.8299999987, 'Elite'],
        ['Adidas', 359, 651661.1200000002, 'Budget']], \
        "Did you correctly calculate values for Adidas products? Expected something different."
    assert last_output_df[4:].values.tolist() == [['Nike', 357, 595341.0199999992, 'Budget'],
        ['Nike', 82, 128475.59000000003, 'Elite'],
        ['Nike', 90, 71843.15000000004, 'Expensive'],
        ['Nike', 16, 6623.5, 'Average']], \
        "Did you correctly calculate values for Nike products? Expected something different."

```

```
['Nike', 16, 6623.5, 'Average']], \
"Did you correctly calculate values for Nike products? Expected something differer
```

Out[7]: 3/3 tests passed

4. Average discount by brand

Interestingly, grouping products by brand and price range allows us to see that Adidas items generate more total revenue regardless of price category! Specifically, "Elite" Adidas products priced \$129 or more typically generate the highest revenue, so the company can potentially increase revenue by shifting their stock to have a larger proportion of these products!

Note we have been looking at `listing_price` so far. The `listing_price` may not be the price that the product is ultimately sold for. To understand `revenue` better, let's take a look at the `discount`, which is the percent reduction in the `listing_price` when the product is actually sold. We would like to know whether there is a difference in the amount of `discount` offered between brands, as this could be influencing `revenue`.

In [8]: `%%sql`

```
-- Select brand and average_discount as a percentage
-- Join brands to finance on product_id
-- Aggregate by brand
-- Filter for products without missing values for brand
```

```
SELECT
    b.brand,
    AVG(f.discount) * 100 AS average_discount

FROM brands AS b
INNER JOIN finance AS f
    ON b.product_id = f.product_id
GROUP BY
    b.brand
HAVING b.brand IS NOT NULL
;
```

```
* postgresql:///sports
2 rows affected.
```

Out[8]:

brand	average_discount
-------	------------------

Nike	0.0
------	-----

Adidas	33.452427184465606
--------	--------------------

In [9]: `%%nose`

```
last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert set(last_output_df.columns) == set(['brand', 'average_discount']), \
```



```

    """Did you select the correct columns? Expected "brand" and "average_discount"."""

def test_shape():
    assert last_output_df.shape[0] == 2, \
        "Did you group by brand? Expected two rows, one per brand."
    assert last_output_df.shape[1] == 2, \
        "Did you select two columns?"

def test_values():
    assert last_output_df.iloc[:, 1].values.tolist() == [0.0, 33.452427184465606], \
        "Did you correctly calculate the average discount for the two brands?"

```

Out[9]: 3/3 tests passed

5. Correlation between revenue and reviews

Strangely, no `discount` is offered on Nike products! In comparison, not only do Adidas products generate the most revenue, but these products are also heavily discounted!

To improve revenue further, the company could try to reduce the amount of discount offered on Adidas products, and monitor sales volume to see if it remains stable. Alternatively, it could try offering a small discount on Nike products. This would reduce average revenue for these products, but may increase revenue overall if there is an increase in the volume of Nike products sold.

Now explore whether relationships exist between the columns in our database. We will check the strength and direction of a correlation between `revenue` and `reviews`.

```

In [10]: %%sql

-- Calculate the correlation between reviews and revenue as review_revenue_corr
-- Join the reviews and finance tables on product_id

SELECT CORR(r.reviews, f.revenue) AS review_revenue_corr
FROM reviews AS r
JOIN finance AS f
    ON r.product_id = f.product_id

* postgresql:///sports
1 rows affected.

```

Out[10]: **review_revenue_corr**
0.6518512283481301

```

In [11]: %%nose

last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert set(last_output_df.columns) == set(['review_revenue_corr']), \
        """Did you calculate the correlation between reviews and revenue, aliasing as "rev

```

```
def test_shape():
    assert last_output_df.shape == (1, 1), \
        "Did you calculate the correlation between reviews and revenue?"

def test_values():
    assert last_output_df.values.tolist() == [[0.6518512283481301]], \
        "Did you correctly calculate how reviews correlates with revenue?"
```

Out[11]: 3/3 tests passed

6. Ratings and reviews by product description length

Interestingly, there is a strong positive correlation between `revenue` and `reviews`. This means, potentially, if we can get more reviews on the company's website, it may increase sales of those items with a larger number of reviews.

Perhaps the length of a product's `description` might influence a product's `rating` and `reviews` — if so, the company can produce content guidelines for listing products on their website and test if this influences `revenue`. Let's check this out!

In [12]: `%%sql`

```
-- Calculate description_length
-- Convert rating to an integer and calculate average_rating
-- Join info to reviews on product_id and group the results by description_length
-- Filter for products without missing values for description, and sort results by description_length

SELECT
    TRUNC(LENGTH(i.description), -2) AS description_length,
    ROUND(AVG(CAST(r.rating AS numeric)), 2) AS average_rating
FROM info AS i
JOIN reviews AS r
    ON i.product_id = r.product_id
WHERE i.description IS NOT NULL
GROUP BY TRUNC(LENGTH(i.description), -2), description_length
ORDER BY description_length;
```

* postgresql:///sports
7 rows affected.

Out[12]: **description_length** **average_rating**

0	1.87
100	3.21
200	3.27
300	3.29
400	3.32
500	3.12
600	3.65

In [13]: `%%nose`

```

last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert set(last_output_df.columns) == set(['description_length', 'average_rating'])
    """Did you select the correct columns use the aliases "description_length" and "av

def test_shape():
    assert last_output_df.shape[0] == 7, \
    """Did you create bins of 100 characters for "description_length"? Expected the ou
    assert last_output_df.shape[1] == 2, \
    "Expected the output to contain two columns."

def test_values():
    last_output_df = last_output.DataFrame().values.astype("float")
    assert last_output_df[0].tolist() == [0.0, 1.87], \
    """Did you sort the results by "description_length" in ascending order?"""
    assert last_output_df[-1].tolist() == [600.0, 3.65], \
    "Did you correctly calculate the results? Expected a different average rating for

```

Out[13]: 3/3 tests passed

7. Reviews by month and brand

Unfortunately, there doesn't appear to be a clear pattern between the length of a product's `description` and its `rating`.

As we know a correlation exists between `reviews` and `revenue`, one approach the company could take is to run experiments with different sales processes encouraging more reviews from customers about their purchases, such as by offering a small discount on future purchases.

Let's take a look at the volume of `reviews` by month to see if there are any trends or gaps we can look to exploit.

```

In [14]: %%sql

-- Select brand, month from last_visited, and a count of all products in reviews alias
-- Join traffic with reviews and brands on product_id
-- Group by brand and month, filtering out missing values for brand and month
-- Order the results by brand and month

SELECT
    b.brand,
    DATE_PART('month', t.last_visited) AS month,
    COUNT(r.*) AS num_reviews
FROM brands AS b
INNER JOIN traffic AS t
    ON b.product_id = t.product_id
INNER JOIN reviews AS r
    ON b.product_id = r.product_id
GROUP BY
    b.brand,
    DATE_PART('month', t.last_visited)
HAVING

```

```

b.brand IS NOT NULL
AND DATE_PART('month', t.last_visited) IS NOT NULL
ORDER BY
b.brand,
DATE_PART('month', t.last_visited)

```

```

* postgresql:///sports
24 rows affected.

```

Out[14]: **brand month num_reviews**

Adidas	1.0	253
Adidas	2.0	272
Adidas	3.0	269
Adidas	4.0	180
Adidas	5.0	172
Adidas	6.0	159
Adidas	7.0	170
Adidas	8.0	189
Adidas	9.0	181
Adidas	10.0	192
Adidas	11.0	150
Adidas	12.0	190
Nike	1.0	52
Nike	2.0	52
Nike	3.0	55
Nike	4.0	42
Nike	5.0	41
Nike	6.0	43
Nike	7.0	37
Nike	8.0	29
Nike	9.0	28
Nike	10.0	47
Nike	11.0	38
Nike	12.0	35

In [15]: **%%nose**

```

last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert set(last_output_df.columns) == set(['brand', 'month', 'num_reviews']), \
        """Did you select the correct columns? Expected "brand", "month", and "num_reviews"

```

```
def test_shape():
    assert last_output_df.shape[0] == 24, \
        "Did you group by brand and month?"
    assert last_output_df.shape[1] == 3, \
        "Did you select three columns?"

def test_values():
    assert last_output_df.iloc[0].values.tolist() == ['Adidas', 1.0, 253], \
        "Expected the first row to contain the number of reviews for Adidas products in Jan"
    assert last_output_df.iloc[-1].values.tolist() == ['Nike', 12.0, 35.0], \
        "Expected the last row to contain the number of reviews for Nike products in Decem"
    assert max(last_output_df["num_reviews"]) == 272, \
        "Did you correctly calculate the number of reviews? Expected the largest number of"
```

Out[15]: 3/3 tests passed

8. Footwear product performance

Looks like product reviews are highest in the first quarter of the calendar year, so there is scope to run experiments aiming to increase the volume of reviews in the other nine months!

So far, we have been primarily analyzing Adidas vs Nike products. Now, let's switch our attention to the type of products being sold. As there are no labels for product type, we will create a Common Table Expression (CTE) that filters `description` for keywords, then use the results to find out how much of the company's stock consists of footwear products and the median `revenue` generated by these items.

```
In [16]: %%sql

-- Create the footwear CTE, containing description and revenue
-- Filter footwear for products with a description containing %shoe%, %trainer, or %fo
-- Also filter for products that are not missing values for description
-- Calculate the number of products and median revenue for footwear products

WITH footwear AS
(SELECT
    i.description,
    f.revenue
FROM info AS i
INNER JOIN finance AS f
    ON i.product_id = f.product_id
WHERE i.description ILIKE '%shoe%'
    OR i.description ILIKE '%trainer%'
    OR i.description ILIKE '%foot%'
    AND i.description IS NOT NULL
)
SELECT
    COUNT(*) AS num_footwear_products,
    PERCENTILE_DISC(0.5) WITHIN GROUP (ORDER BY revenue) AS median_footwear_revenue
FROM footwear

* postgresql:///sports
1 rows affected.
```

Out[16]: **num_footwear_products** **median_footwear_revenue**

2700	3118.36
------	---------

In [17]: **%%nose**

```
last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert set(last_output_df.columns) == set(['num_footwear_products', 'median_footwear_revenue'])
    "Did you select the correct columns and use the aliases specified in the instructions?"

def test_shape():
    assert last_output_df.shape[0] == 1, \
        "Expected the output to contain one row."
    assert last_output_df.shape[1] == 2, \
        "Expected the output to contain two columns."

def test_values():
    assert last_output_df.iloc[0,0] == 2700, \
        "Did you count the number of footwear products?"
    assert last_output_df.iloc[0,1] == 3118.36, \
        "Did you calculate the median revenue for footwear products?"
```

Out[17]: 3/3 tests passed

9. Clothing product performance

Recall from the first task that we found there are 3,117 products without missing values for `description`. Of those, 2,700 are footwear products, which accounts for around 85% of the company's stock. They also generate a median revenue of over \$3000 dollars!

This is interesting, but we have no point of reference for whether footwear's `median_revenue` is good or bad compared to other products. So, for our final task, let's examine how this differs to clothing products. We will re-use `footwear`, adding a filter afterward to count the number of products and `median_revenue` of products that are not in `footwear`.

In [18]: **%%sql**

```
-- Copy the footwear CTE from the previous task
-- Calculate the number of products in info and median revenue from finance
-- Inner join info with finance on product_id
-- Filter the selection for products with a description not in footwear

WITH footwear AS
(SELECT
    i.description,
    f.revenue
FROM info AS i
INNER JOIN finance AS f
    ON i.product_id = f.product_id
WHERE i.description ILIKE '%shoe%'
    OR i.description ILIKE '%trainer%')
```

```

        OR i.description ILIKE '%foot%'
        AND i.description IS NOT NULL
    )

    SELECT
        COUNT(i.*) AS num_clothing_products,
        PERCENTILE_DISC(0.5) WITHIN GROUP (ORDER BY f.revenue) AS median_clothing_revenue
    FROM info AS i
    INNER JOIN finance AS f
        ON i.product_id = f.product_id
    WHERE i.description NOT IN
        (SELECT DISTINCT description
         FROM footwear)

```

* postgresql:///sports
1 rows affected.

Out[18]: **num_clothing_products** **median_clothing_revenue**

417	503.82
-----	--------

In [19]: **%%nose**

```

last_output = _
last_output_df = last_output.DataFrame()

def test_columns():
    assert set(last_output_df.columns) == set(['num_clothing_products', 'median_clothing_revenue'])
    "Did you select the correct columns and use the aliases specified in the instructions?"

def test_shape():
    assert last_output_df.shape[0] == 1, \
        "Expected the output to contain one row."
    assert last_output_df.shape[1] == 2, \
        "Expected the output to contain two columns."

def test_values():
    assert last_output_df.iloc[0,0] == 417, \
        "Did you count the number of clothing products? Expected there to be 417 items."
    assert last_output_df.iloc[0,1] == 503.82, \
        "Did you calculate the median revenue for clothing products? Expected it to be $503.82"

```

Out[19]: 3/3 tests passed